《人工智能基础》实验报告三

2025 年 4 月 1 日

0.1 Package import

```python
[25]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split,GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.metrics import accuracy_score, precision_score, recall_score,␣
 ↪f1_score, confusion_matrix, roc_curve,␣
 ↪auc,precision_recall_curve,average_precision_score


import numpy as np
```

0.2 Dataset preprocessing and splitting

```python
[26]: data = load_breast_cancer()
X = data.data
y = data.target


scaler = StandardScaler()
X = scaler.fit_transform(X)


X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,␣
 ↪random_state=3149)
```

0.3 Grid Search for hyperparameter tuning

```python
param_grid = {
    'C': np.linspace(0.5, 1.5, 10),  # 正则化
    'penalty': ['l1', 'l2'],
    'solver': ["liblinear","saga"],
    'max_iter': np.arange(50, 150, 10),  # 迭代次数
    'class_weight': [None, 'balanced'],
    'fit_intercept': [True, False],
    'intercept_scaling': [1, 2],
    'multi_class': ['auto', 'ovr'],
}

grid_search = GridSearchCV(LogisticRegression(), param_grid, cv=10,
    scoring='accuracy')
grid_search.fit(X_train, y_train)

print("最佳参数:", grid_search.best_params_)
print("最佳得分:", grid_search.best_score_)
print("最佳模型:", grid_search.best_estimator_)
```

0.4 Train and test the model

```python
model = LogisticRegression(C=np.float64(0.8),
                           fit_intercept=False,
                           max_iter=np.int64(110),
                           multi_class='auto',
                           penalty='l1',
                           solver='saga'
                           )
model.fit(X_train, y_train)



y_pred = model.predict(X_test)
y_pred_proba = model.predict_proba(X_test)[:, 1]
```

/home/yee/conda/envs/zju-yee/lib/python3.10/site-
packages/sklearn/linear_model/_logistic.py:1247: FutureWarning: 'multi_class'

was deprecated in version 1.5 and will be removed in 1.7. From then on, it will always use 'multinomial'. Leave it to its default value to avoid this warning.
  warnings.warn(
/home/yee/conda/envs/zju-yee/lib/python3.10/site-packages/sklearn/linear_model/_sag.py:348: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
  warnings.warn(

0.5 Model evaluation

```python
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print(f'准确率 (Accuracy): {accuracy:.2f}')
print(f'精确率 (Precision): {precision:.2f}')
print(f'召回率 (Recall): {recall:.2f}')
print(f'F1 分数 (F1 Score): {f1:.2f}')




conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 6))
plt.imshow(conf_matrix, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion Matrix')
plt.colorbar()
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()




fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = auc(fpr, tpr)


plt.figure(figsize=(8, 6))
```

```python
plt.plot(fpr, tpr, label=f'ROC curve (area = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()
plt.grid(True)
plt.show()




precision, recall, _ = precision_recall_curve(y_test, y_pred_proba)
average_precision = average_precision_score(y_test, y_pred_proba)
plt.figure(figsize=(8, 6))
plt.plot(recall, precision, label=f'PR curve (area = {average_precision:.2f})')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.legend()
plt.grid(True)
plt.show()
```

准确率 (Accuracy): 0.97
精确率 (Precision): 0.96
召回率 (Recall): 0.99
F1 分数 (F1 Score): 0.97

Confusion Matrix

Receiver Operating Characteristic (ROC) Curve

ROC curve (area = 1.00)

True Positive Rate

False Positive Rate

Precision-Recall Curve